

Homework #4

Mahan Fathi

March 30, 2022

1 Question 1




Figure 1: **Q1** MsPacman-v0 Single DQN

```
python run_hw4.py exp_name=q1 env_name=MsPacman-v0 n_iter=10000000
```

Listing 1: **Q1** Run command

2 Question 2

The plummet after reaching a total return of 150 has already been alluded to in the problem statement, and is the case also here.

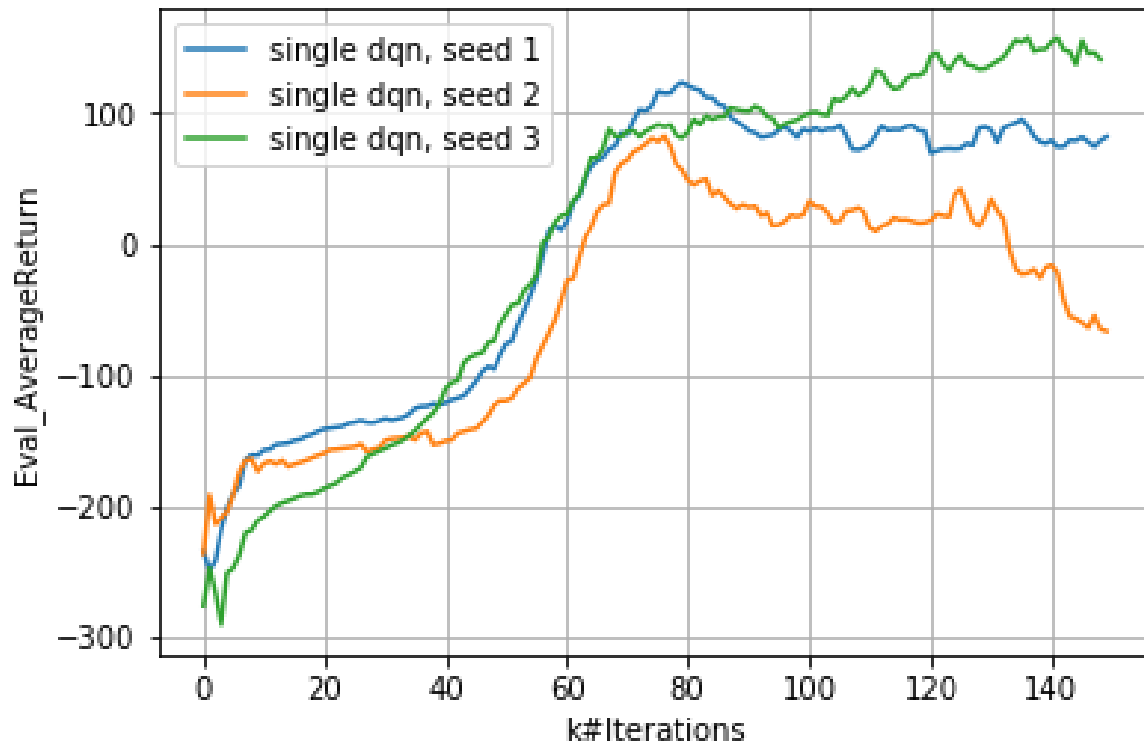


Figure 2: **Q2** LunarLander-v3 Single DQN

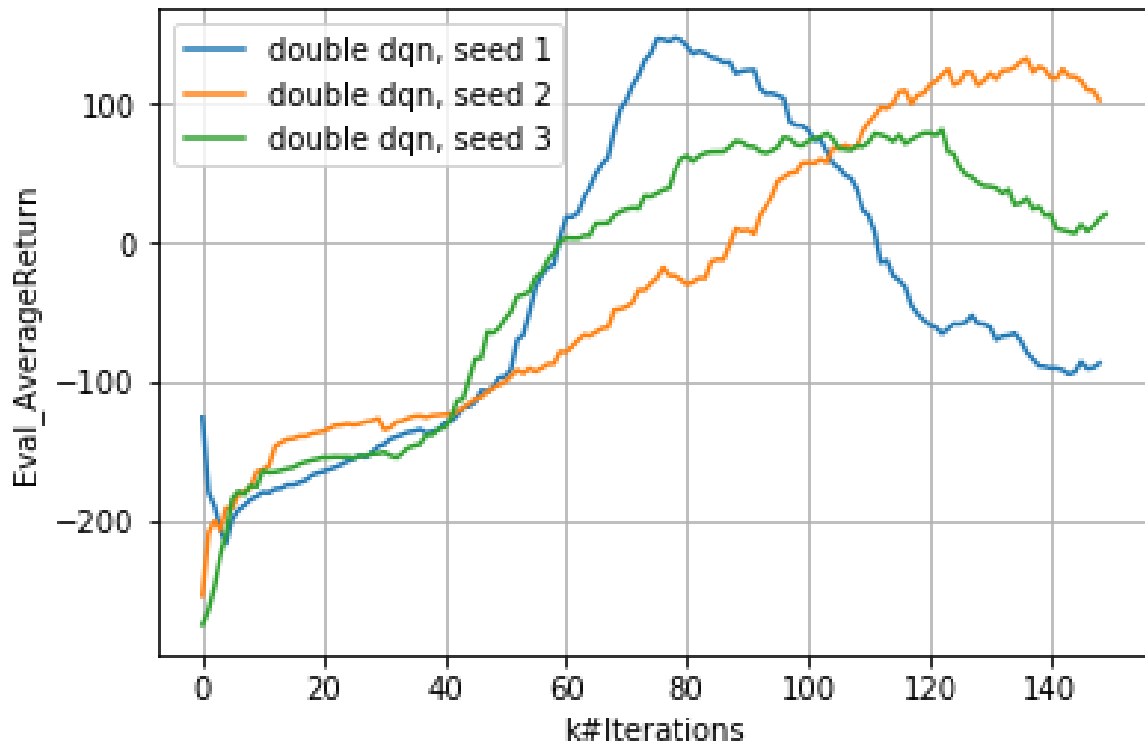


Figure 3: **Q2** LunarLander-v3 Double DQN

```
python run_hw4.py env_name=LunarLander-v3 \
    exp_name=q2_dqn_1 seed=1 n_iter=1000000 double_q=null

python run_hw4.py env_name=LunarLander-v3 \
    exp_name=q2_dqn_2 seed=2 n_iter=1000000 double_q=null

python run_hw4.py env_name=LunarLander-v3 \
    exp_name=q2_dqn_3 seed=3 n_iter=1000000 double_q=null

python run_hw4.py env_name=LunarLander-v3 \
    exp_name=q2_doubledqn_1 seed=1 n_iter=200000

python run_hw4.py env_name=LunarLander-v3 \
    exp_name=q2_doubledqn_2 seed=2 n_iter=200000

python run_hw4.py env_name=LunarLander-v3 \
    exp_name=q2_doubledqn_3 seed=3 n_iter=200000
```

Listing 2: **Q2** Run commands

3 Question 3




Figure 4: Q3 LunarLander-v3 with ablations for gamma

```
python run_hw4.py env_name=LunarLander-v3 \  
    exp_name=q3_gamma0.9 n_iter=1000000 gamma=0.9  
  
python run_hw4.py env_name=LunarLander-v3 \  
    exp_name=q3_gamma0.95 n_iter=1000000 gamma=0.95  
  
python run_hw4.py env_name=LunarLander-v3 \  
    exp_name=q3_gamma0.99 n_iter=1000000 gamma=0.99
```

Listing 3: Q3 Run commands

Here we have chosen `params['gamma']` as our ablation variable. The reason behind this is that one convenient handle we have over the bias and variance trade-off in a Q-Learning algorithm, is the discount factor. My speculation was that the fall in the return after 150 could be addressed by making the Q values more or less sensitive to the changes in the rewarding pattern of the environment. We can see from plots that indeed $\gamma = 0.99$ performs best, as it corresponds to a scenario in which the algorithm is least fixated on the current reward and takes into account the future rewards more, as opposed to other γ values.

4 Question 4

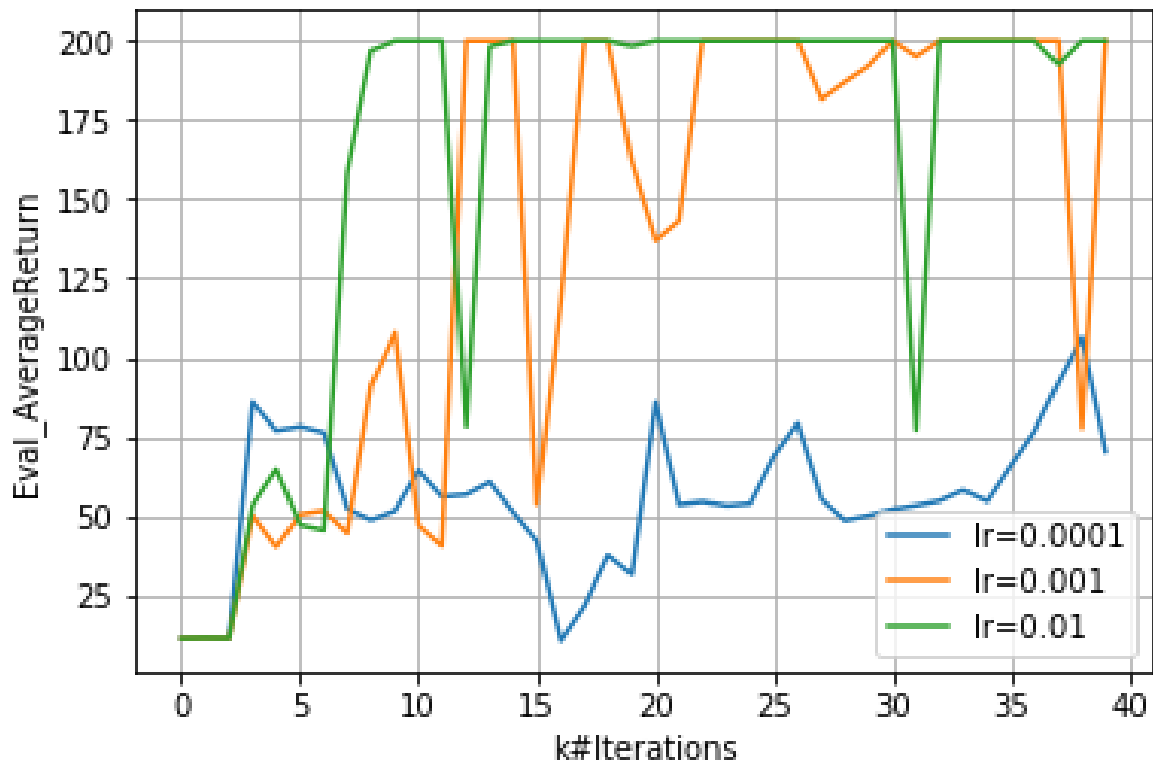


Figure 5: Q4 DDPG for InvertedPendulum-v2, learning_rate ablations

```
python run_hw4.py exp_name=q4_ddpg_lr0.001 rl_alg=ddpg \
    env_name=InvertedPendulum-v2 atari=false n_iter=40000 learning_rate=0.001

python run_hw4.py exp_name=q4_ddpg_lr0.01 rl_alg=ddpg \
    env_name=InvertedPendulum-v2 atari=false n_iter=40000 learning_rate=0.01

python run_hw4.py exp_name=q4_ddpg_lr0.0001 rl_alg=ddpg \
    env_name=InvertedPendulum-v2 atari=false n_iter=40000 learning_rate=0.0001

python run_hw4.py exp_name=q4_ddpg_uf1 rl_alg=ddpg \
    env_name=InvertedPendulum-v2 atari=false n_iter=40000 policy_update_frequency=1

python run_hw4.py exp_name=q4_ddpg_uf10 rl_alg=ddpg \
    env_name=InvertedPendulum-v2 atari=false n_iter=40000 policy_update_frequency=10

python run_hw4.py exp_name=q4_ddpg_uf100 rl_alg=ddpg \
    env_name=InvertedPendulum-v2 atari=false n_iter=40000 policy_update_frequency=100
```

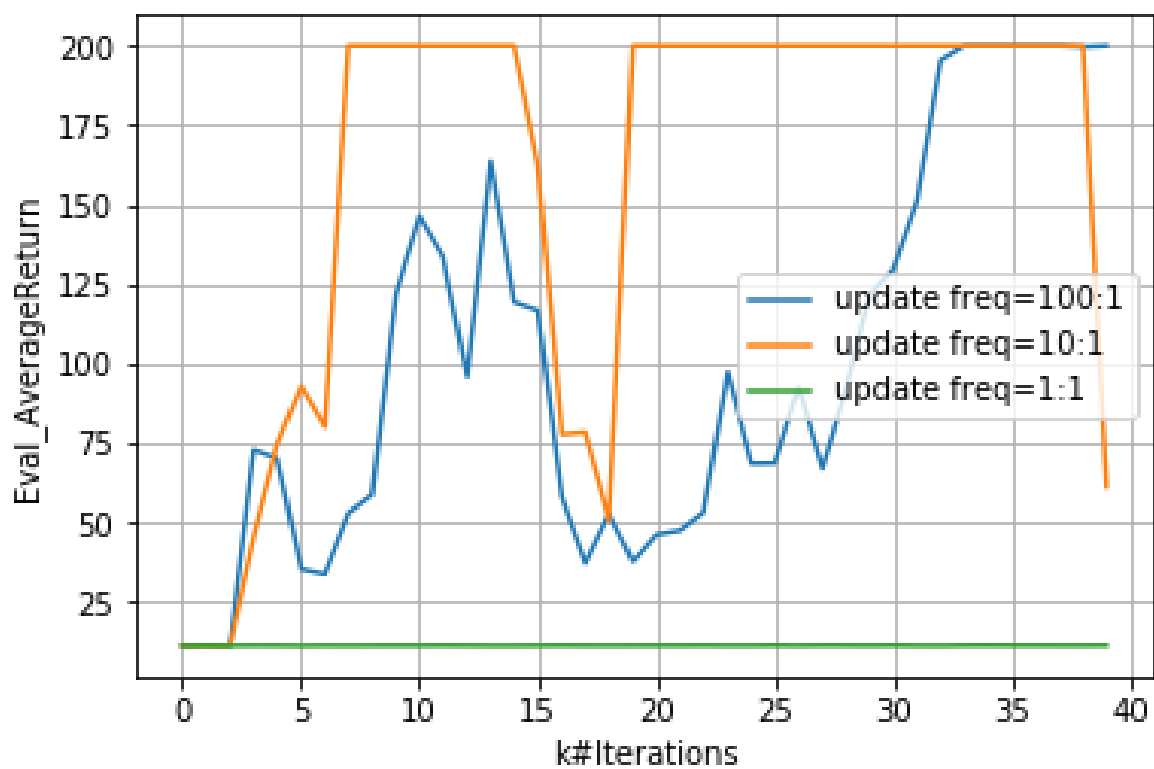


Figure 6: **Q4** DDPG for InvertedPendulum-v2, policy_update_frequency ablations

5 Question 5

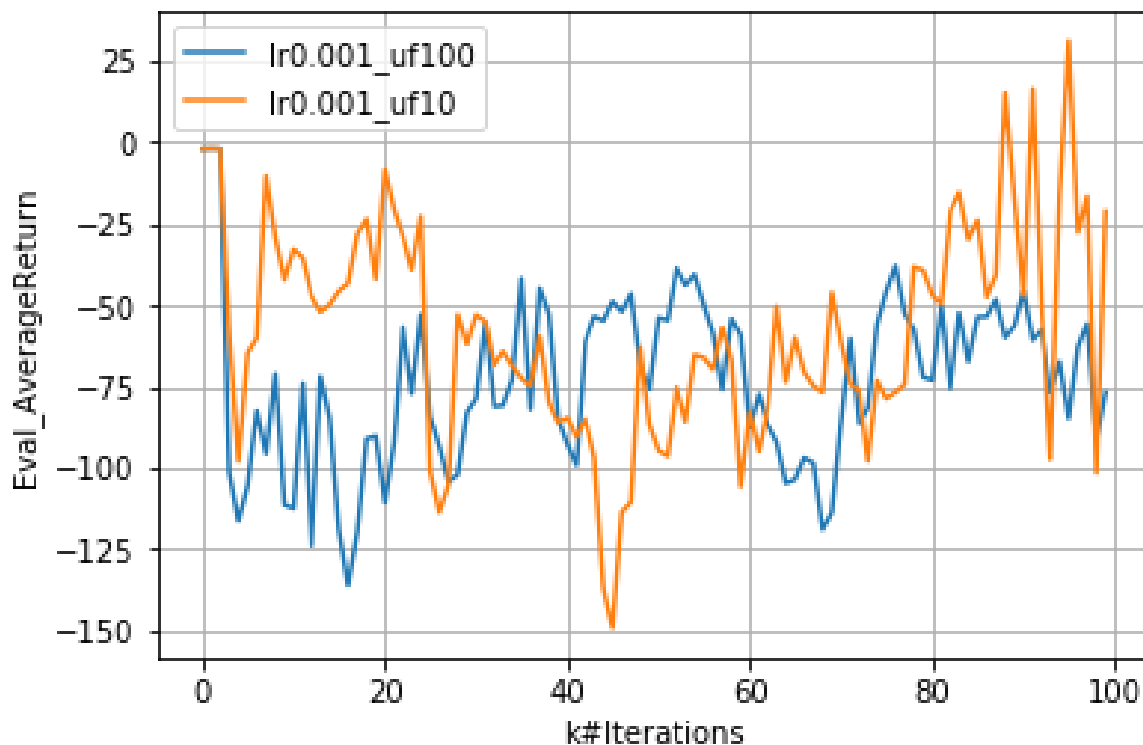


Figure 7: **Q5** DDPG for HalfCheetah-v2, with two different learning schedule

DDPG without the exploration noise during data collection does not seem to be working readily out-of-the-box for the HalfCheetah-v2 environment, however, the good news is that TD3 is successful in doing so. Take a look at Question 7.

```
python run_hw4.py exp_name=q5_ddpg_hard_lr0.001_uf10 rl_alg=ddpg \
  env_name=HalfCheetah-v2 atari=false n_iter=100000 \
  learning_rate=0.001 policy_update_frequency=10
```

```
python run_hw4.py exp_name=q5_ddpg_hard_lr0.001_uf100 rl_alg=ddpg \
  env_name=HalfCheetah-v2 atari=false n_iter=100000 \
  learning_rate=0.001 policy_update_frequency=100
```


6 Question 6

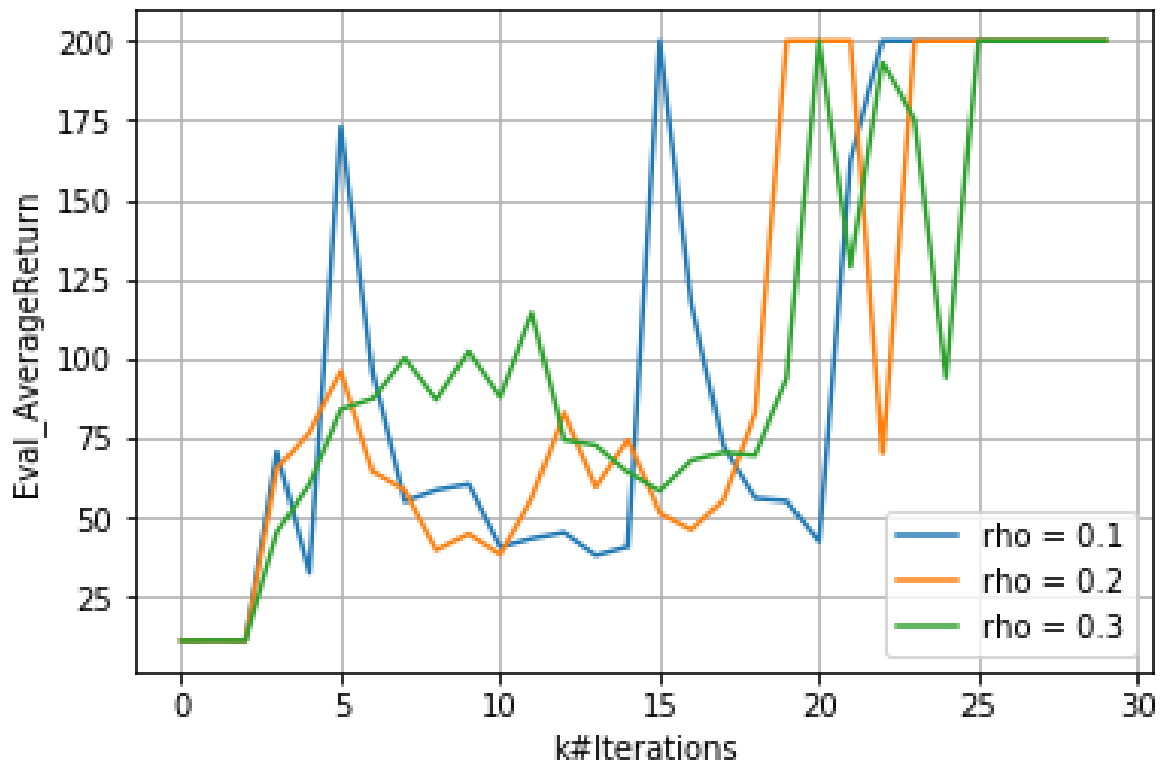


Figure 8: Q6 TD3 for InvertedPendulum-v2, td3_target_policy_noise ablations

```
python run_hw4.py exp_name=q6_td3_rho0.1 rl_alg=td3 \
  env_name=InvertedPendulum-v2 atari=false n_iter=30000 \
  learning_rate=0.001 td3_target_policy_noise=0.1

python run_hw4.py exp_name=q6_td3_rho0.2 rl_alg=td3 \
  env_name=InvertedPendulum-v2 atari=false n_iter=30000 \
  learning_rate=0.001 td3_target_policy_noise=0.2

python run_hw4.py exp_name=q6_td3_rho0.3 rl_alg=td3 \
  env_name=InvertedPendulum-v2 atari=false n_iter=30000 \
  learning_rate=0.001 td3_target_policy_noise=0.3

python run_hw4.py exp_name=q6_td3_rho0.1_shape2 rl_alg=td3 \
  env_name=InvertedPendulum-v2 atari=false n_iter=100000 \
  learning_rate=0.001 td3_target_policy_noise=0.1 n_layers_critic=2
```

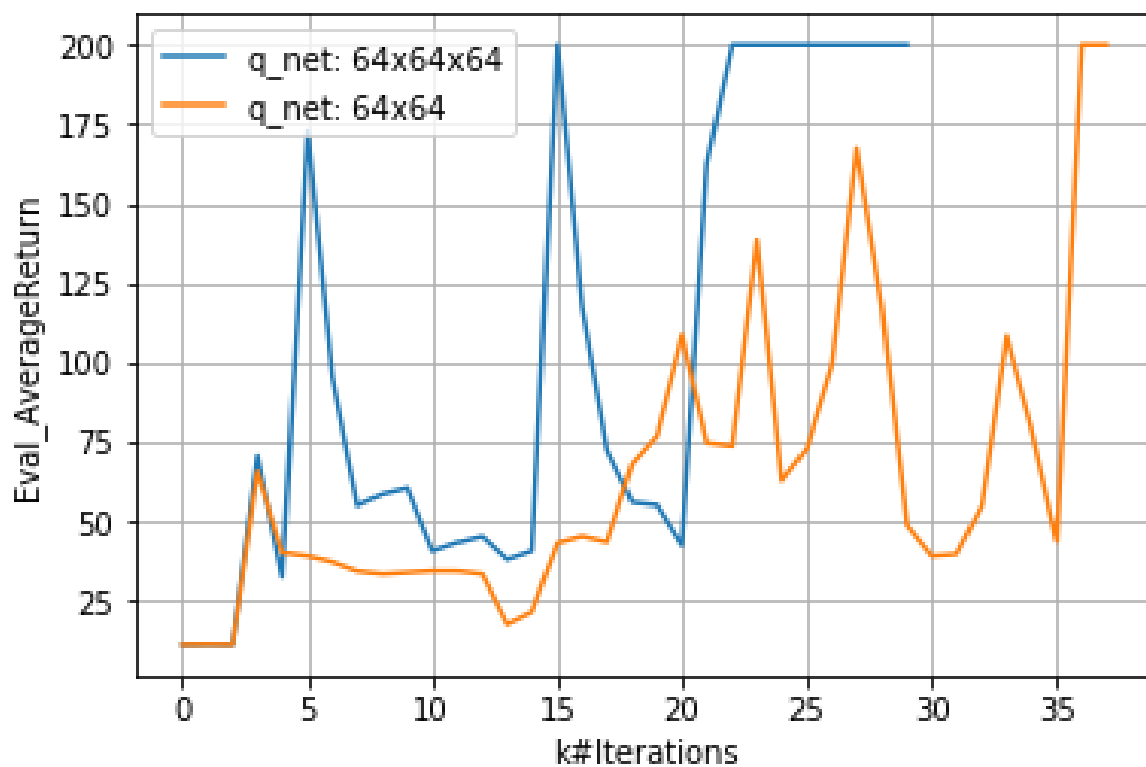


Figure 9: **Q6** TD3 for InvertedPendulum-v2, Q-network size ablations

7 Question 7

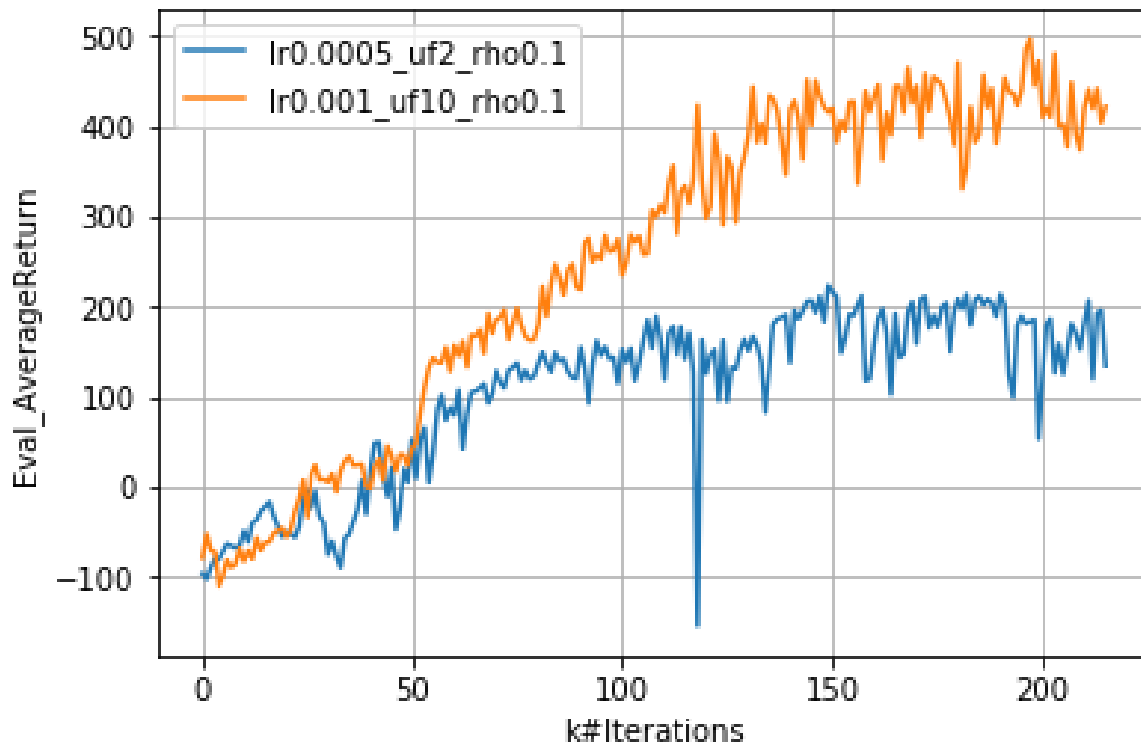


Figure 10: **Q7** TD3 for HalfCheetah-v2

```
python run_hw4.py exp_name=q7_td3_hard_lr0.001_uf10_rho0.1_shape3 \  
  rl_alg=td3 n_layers_critic=3 \  
  env_name=HalfCheetah-v2 atari=false n_iter=1000000 \  
  learning_rate=0.001 policy_update_frequency=10 td3_target_policy_noise=0.1
```

```
python run_hw4.py exp_name=q7_td3_hard_lr0.0005_uf2_rho0.1_shape3 \  
  rl_alg=td3 n_layers_critic=3 \  
  env_name=HalfCheetah-v2 atari=false n_iter=1000000 \  
  learning_rate=0.0005 policy_update_frequency=2 td3_target_policy_noise=0.1
```